

OraSRS: A Compliant and Lightweight Decentralized Threat Intelligence Protocol with Time-Bounded Risk Enforcement

LuoZiQian
{luo.zi.qian@orasrs.net}

December 14, 2025

Abstract

I present OraSRS, a protocol security and compliance engineering framework that addresses a fundamental tension in security systems: the trade-off between response speed and decentralized trust. Rather than focusing on blockchain as a core innovation, OraSRS leverages blockchain as an infrastructure layer to coordinate trust in a protocol-safe manner while prioritizing $\leq 100\text{ms}$ response for real-time defense. After implementing and testing several centralized and decentralized approaches, I observed that existing solutions inadequately balance speed, compliance, and security. OraSRS achieves sub-100ms local response through optimistic execution at the edge while maintaining global consistency via protocol-based coordination. Key innovations include: (1) a T0-T3 verification lifecycle decoupling local defense from final confirmation; (2) a commit-reveal protocol with economic incentives deterring Sybil attacks; (3) time-bounded risk enforcement ensuring regulatory compliance. Unlike pure blockchain protocols, OraSRS emphasizes protocol engineering for real-world deployability and compliance. Evaluation demonstrates sub-0.04ms local latency, 95.5% F1-score, and $\geq 95\%$ Sybil defense rate. The system is open-sourced at <https://github.com/srs-protocol/OraSRS-protocol/> under Apache 2.0.

Keywords: Protocol security engineering, Threat intelligence coordination, Optimistic verification, Compliance-aware security, Sybil attack prevention

1 Introduction

1.1 Research Background

As an independent security researcher, I have observed that traditional threat intelligence services face a fundamental dilemma: centralized solutions provide fast response times but create single points of failure, while decentralized systems offer security at the cost of latency that is simply unacceptable for real-time defense.

The core challenge lies not in the blockchain technology itself, but in protocol engineering that balances speed, security, and compliance. Unlike corporate or academic systems developed by large teams, OraSRS is built by a single developer to address real-world operational gaps that are often overlooked by institutional projects.

In my experience developing security systems, I found that existing decentralized threat intelligence systems often overlook the critical need for real-time defense capabilities, resulting in response times of 100-500ms due to consensus mechanisms. These approaches fail to meet two critical needs in modern networks: (a) sub-100ms response for real-time defense, and (b) auditability without central control.

While blockchain-based TI systems improve transparency, beginning with the foundational work of Nakamoto on Bitcoin [1], they introduce unacceptable latency (100-500ms due to consensus mechanisms), making them unsuitable for real-time threat blocking. To bridge this gap, I designed OraSRS as a protocol engineering solution that leverages blockchain as a trust coordination infrastructure while prioritizing response speed through optimistic verification.

This design separates risk assessment from enforcement: clients receive consultative scores and decide locally whether to block, enabling rapid threat response while maintaining security benefits.

The challenge I encountered was that pure blockchain-centric approaches, while providing auditability, introduce delays (100-500ms due to consensus mechanisms) that make them unsuitable for immediate threat blocking. For example, when my system detects an IP address launching a DDoS attack, waiting for blockchain confirmation means the attack may have already caused damage. This timing issue highlighted the need for protocol-level solutions that decouple immediate response from global consensus, which is the core engineering innovation of OraSRS.

1.2 Research Objectives and Contributions

My work addresses the fundamental tension between speed and security in threat intelligence sharing. The main objectives include:

1. **Designing a local-first threat intelligence protocol:** I implemented a decentralized threat intelligence sharing network that achieves sub-100ms response while maintaining blockchain-level security guarantees.
2. **Implementing consultative risk assessment:** I adopted a consultative rather than blocking approach, providing risk ratings and recommendations to network entities while leaving final blocking decisions to clients.
3. **Ensuring privacy and compliance:** I designed threat intelligence sharing while strictly protecting user privacy and meeting data protection requirements of various jurisdictions.
4. **Achieving auditability and transparency:** I used decentralized consensus technology for tamper-proof records and complete audit trails of threat intelligence.
5. **Improving threat detection accuracy:** I enhanced threat detection accuracy and timeliness through distributed consensus mechanisms and commit-reveal protocols.

Through my implementation, I found that the OraSRS protocol effectively addresses the speed vs. security trade-off through protocol security engineering. My approach emphasizes practical protocol design over pure blockchain innovation, focusing on security engineering principles. The main contributions include:

- A novel protocol security engineering approach with optimistic verification architecture and T0-T3 time model that achieves $\leq 100\text{ms}$ threat response while maintaining decentralized security - this represents my biggest protocol engineering innovation.
- A commit-reveal mechanism with staking and slashing as a protocol-level defense against Sybil attacks and lazy validators, with formal economic security analysis showing Nash equilibrium properties.
- A compliance-aware protocol design using Chinese SM algorithms to meet security compliance requirements of specific regions, demonstrating protocol-to-regulation mapping.
- A performance evaluation framework for protocol security systems, demonstrating 0.03ms local latency, 95.5% F1-score, and $\geq 95\%$ Sybil defense rate as measurable security engineering outcomes.

2 Related Work

2.1 Traditional Threat Intelligence Services

Traditional threat intelligence services like VirusTotal and IBM X-Force provide centralized threat intelligence query services. While these services offer fast response times (typically $\leq 50\text{ms}$) and provide public APIs that offer a degree of transparency, I found they suffer from single point of failure risks, data bias, and privacy leakage issues. They rely on centralized trust mechanisms and typically use blocking methods, directly blocking network traffic without full transparency. My experience indicates that these centralized solutions are vulnerable to targeted attacks and regulatory pressures. Prior work in this area often overlooks the temporal nature of threats and the need for immediate response capabilities.

2.2 Decentralized Threat Intelligence Sharing

Recent developments in decentralized threat intelligence include several frameworks that attempt to address the limitations of centralized systems. For example, ThreatExchange proposed by Facebook allows multiple organizations to share threat intelligence but still relies on centralized coordination mechanisms. CIF (Collective Intelligence Framework) provides a standardized threat intelligence format and sharing protocol but still has deficiencies in decentralization and trust mechanisms. The STIX/TAXII standard defines structured representation methods and transmission protocols for threat intelligence but also lacks decentralized trust mechanisms. In my evaluation of these systems, I observed that they often trade performance for decentralization without achieving a practical balance for real-world deployment.

Table 1 provides a detailed comparison of these existing systems with OraSRS, highlighting the unique advantages of our approach:

System	Latency	Decentralization	Compliance	Open	Key Li
VirusTotal	>50ms	Low	Low	Yes (API)	Relies on ce
CIF	50-100ms	Medium	Medium	Yes	Limited trus
STIX/TAXII	100-200ms	Medium	High	Yes	Complex im
ThreatExchange	>100ms	Low	Medium	Yes	Centralized
OraSRS (Proposed)	<100ms	High	High	Yes	Introduces token e

Table 1: Comparison of Threat Intelligence Systems

2.3 Decentralized Consensus Layer Threat Intelligence

I investigated several works that have explored blockchain-based threat intelligence sharing. However, existing solutions focus primarily on audit trails without addressing response latency. Through my implementation and testing, I found that pure blockchain approaches result in 100-500ms response times, which is insufficient for real-time defense. This latency issue arises from consensus mechanisms that require multiple network round-trips before finality.

More recent developments in blockchain-based threat intelligence include systems that leverage smart contracts for automated threat response, though many lack the optimistic verification mechanism that enables OraSRS’s sub-100ms response [2]. Some frameworks provide decentralized threat detection but require significant computational resources and lack consensus mechanisms for global coordination [8]. Other systems address privacy concerns but introduce additional latency due to encryption operations [7].

Furthermore, some approaches provide strong privacy guarantees but are limited to small-scale deployments due to communication complexity [3]. Other systems ensure privacy while maintaining data integrity, yet require complex cryptographic setup and verification processes that impact real-time performance [7]. Some frameworks combine cloud efficiency with blockchain security but still face challenges in achieving the speed-security balance that OraSRS addresses through its T0-T3 verification lifecycle [5]. Related work in hybrid consensus mechanisms, such as Proof of Activity which extends Bitcoin’s Proof of Work via Proof of Stake, also demonstrates the potential of combining multiple verification approaches [4]. Additionally, foundational work in Ethereum’s implementation has provided insights into decentralized application design [6].

In contrast, OraSRS achieves sub-100ms response through optimistic execution, which I designed specifically to address this performance limitation in existing work. Additionally, OraSRS incorporates Chinese cryptographic algorithms (SM2/SM3/SM4) to meet regional compliance requirements and enhance data security, supporting deployment on national-compliant chains such as ChainMaker.

2.4 Optimistic Verification Mechanisms

Optimistic verification is a mechanism balancing efficiency and security in decentralized systems. Early Plasma frameworks introduced optimistic assumptions, assuming all operations were valid unless someone challenged them. Optimistic Rollups further developed this concept, allowing fast transaction confirmation while retaining challenge periods to ensure security. However, these mechanisms were mainly applied to financial transactions and their application in threat intelligence sharing is first of its kind. I observed that op-

timistic verification is particularly well-suited for threat intelligence due to the temporal nature of threats - most threats are time-sensitive and require immediate action, even if final verification occurs later.

2.5 Commit-Reveal Schemes

Commit-Reveal is a cryptographic protocol widely used to prevent front-running and ensure fairness. In decentralized gambling, auctions, and voting systems, Commit-Reveal mechanisms work by submitting hash values first and revealing original values later, preventing participants from changing strategies after seeing others' choices. In threat intelligence, this represents the first application of such mechanisms to prevent malicious participants from manipulating systems. My approach extends these mechanisms with economic incentives for security, addressing a limitation I identified in prior work where pure cryptographic approaches lacked sufficient economic disincentives for malicious behavior.

3 System Model and Architecture

3.1 Formal Definition

We define the OraSRS system as a seven-tuple (N, S, T, R, P, V, C) , where:

- N : Set of participating nodes
- S : Set of threat intelligence states
- T : Time parameter set, including $T_{detect}, T_{local}, T_{consensus}$
- R : Reward distribution function
- P : Penalty enforcement function
- V : Set of verification mechanisms
- C : Consensus protocol

The OraSRS architecture implements a three-tier consensus architecture: Global Root Network Layer, Partition Consensus Layer, and Edge Caching Layer. This design ensures global synchronization while maintaining local responsiveness. The system incorporates Chinese cryptographic algorithms (SM2/SM3/SM4) to meet regional compliance requirements, particularly for deployments in China that must adhere to Cybersecurity Protection Level 2.0 standards.

3.2 Optimistic Verification Lifecycle

OraSRS's core innovation is the optimistic verification model, combining T0 local defense with T3 global consensus to solve the contradiction between blockchain confirmation delay and security response speed. Our implementation suggests that this approach effectively addresses the fundamental tension between speed and security in threat intelligence systems.

3.2.1 Time Parameter Definition

- T_{detect} : Threat detection time, typically $\approx 10\text{ms}$
- T_{local} : Local activation time, using ipset to achieve $O(1)$ queries, $\approx 1\text{ms}$
- $T_{consensus}$: Global consensus time, dependent on underlying blockchain, typically $\approx 30\text{s}$

3.2.2 Optimistic Verification Process

OraSRS adopts an optimistic verification lifecycle including the following stages:

1. **Local Optimistic Execution (T0)**: Edge nodes immediately execute defensive measures locally after detecting threats
2. **Submit Consensus (T1-T2)**: Submit threat intelligence to decentralized consensus network for verification
3. **Global Confirmation (T3)**: Complete decentralized consensus to form final state
4. **State Synchronization**: Synchronize final state to all nodes

3.3 Triad Architecture Design

I implement OraSRS as a triad architecture with specific design choices based on performance and deployment requirements:

1. **Edge Agent**: I chose kernel-level ipset (not userspace proxy) to achieve $O(1)$ lookup with $\approx 5\text{MB}$ memory. Ultra-lightweight threat detection agents responsible for local threat detection and Local Optimistic State execution. After testing various approaches, I determined that P2P solutions provide limited benefits for data synchronization in this context, so I removed complex P2P schemes in favor of direct client-to-blockchain connections.
2. **Consensus Layer**: I integrate ChainMaker with SM2/3/4 to comply with Chinese regulations. Multi-chain trusted storage ensuring immutability of threat intelligence. This layer implements BFT consensus algorithms where client-uploaded threat information is verified through multi-node consensus and whitelist filtering. For enhanced security against forgery, I designed hardware key signing for threat intelligence uploads.
3. **Intelligence Layer**: Threat intelligence coordination network implementing global threat intelligence aggregation and distribution. The protocol supports ChainMaker deployment for domestic compliance, with data mirroring capabilities to ensure data does not leave national borders.

Edge Layer	Consensus Layer	Intelligence Layer
<i>Local-first enforcement</i>	<i>Tamper-proof verification</i>	<i>Global reputation sharing</i>
<ul style="list-style-type: none"> • Ultra-lightweight agent (< 5MB) • Kernel-level ipset ($O(1)$ lookup) • Local optimistic state • < 100ms response 	<ul style="list-style-type: none"> • BFT consensus with SM2/3/4 • Commit-Reveal mechanism • Staking & slashing • Global finalized state 	<ul style="list-style-type: none"> • Threat intelligence coordination • Global aggregation & distribution • P2P/RPC network • AI-enhanced analysis

Figure 1: OraSRS system architecture showing the triad design: (1) Edge Agent for local-first enforcement, (2) Consensus Layer for tamper-proof verification, and (3) Intelligence Layer for global reputation sharing.

3.4 Decentralized Architecture and Implementation

3.5 Architectural Trade-offs: RPC over P2P

In designing the network topology for OraSRS, we evaluated both pure Peer-to-Peer (P2P) gossip protocols and Remote Procedure Call (RPC) based direct connections. While P2P networks offer theoretical decentralization, we explicitly selected a direct RPC connection model with internal gateway nodes for the following critical engineering and compliance reasons:

Regulatory Compliance and Auditability: Uncontrolled P2P traffic propagation can pose significant compliance risks in strictly regulated network environments (e.g., meeting China’s Cybersecurity Law requirements). P2P protocols often involve indiscriminate data exchange with unknown peers, which complicates traffic auditing and source tracing. By utilizing trusted RPC nodes, the protocol ensures that threat intelligence synchronization remains within controllable, auditable channels, satisfying the “Data Minimization” and sovereignty requirements described in Section 8.

Enterprise Intranet Penetration: Corporate security policies typically strictly block non-standard P2P traffic (e.g., BitTorrent-like protocols) at the firewall level. Implementing a P2P mesh often requires complex NAT traversal techniques (STUN/TURN) which introduces additional latency and security holes. OraSRS utilizes Internal RPC Nodes as cross-chain synchronization gateways. These gateways act as trusted bridges that pull finalized states from the consensus layer and distribute them to internal edge agents via standard HTTP/JSON-RPC protocols, ensuring seamless integration with existing enterprise network topologies without requiring firewall reconfiguration.

Client Complexity and Performance: Maintaining a full P2P routing table consumes significant memory and CPU resources, which contradicts our design goal of an “ultra-lightweight agent” (<5MB). The RPC approach offloads the heavy lifting of consensus synchronization to the node layer, allowing edge agents to focus purely on high-performance packet filtering ($O(1)$ ipset lookup).

3.6 Optimistic Verification Architecture

3.6.1 Local Optimistic State and Global Finalized State

OraSRS maintains two states:

- **Local Optimistic State:** Stored in local ipset for fast queries and blocking
- **Global Finalized State:** Stored on blockchain with immutability and global consistency

3.6.2 Sequence Timeline Description

Here is the detailed sequence flow of OraSRS optimistic verification, solving the contradiction between blockchain confirmation delay and security response speed:

Stage 1: Threat Detection (T0)

- Edge nodes detect malicious IP (e.g., 1.2.3.4)
- Immediately execute defensive measures locally ($\pm 1\text{ms}$)
- Prepare to submit threat intelligence to blockchain simultaneously

Stage 2: Commit Phase (T1)

- Calculate threat intelligence hash: $h = \text{Hash}(IP||threat_level||salt)$
- Submit hash value to decentralized consensus layer (preventing front-running)
- Set commit deadline B_{commit}

Stage 3: Optimistic Execution (T2)

- Local ipset immediately updates to block the IP
- Other network nodes synchronize Local Optimistic State via RPC
- Achieve $\pm 100\text{ms}$ threat response time

Stage 4: Reveal Phase (T3)

- Reveal original threat intelligence after preset time
- Verify $\text{Hash}(IP||threat_level||salt) == h$
- Complete on-chain consensus verification

Stage 5: State Confirmation (T4)

- Verification passes: Threat intelligence written to final state
- Verification fails: Revert Local Optimistic State update
- Execute incentive/punishment mechanisms

Our implementation demonstrates that this design allows the system to achieve near-local defense response speed while maintaining decentralized security, which is OraSRS's biggest architectural innovation.

4 Core Mechanisms

4.1 Risk Scoring Algorithm

We implement a multi-dimensional risk scoring algorithm in OraSRS:

$$RiskScore = \sum_{i=1}^n (weight_i \times timeDecay_i \times sourceMultiplier_i) \quad (1)$$

Where:

- $weight_i$: Weight of threat category i (calibrated based on threat severity statistics from CIC-IDS2017 dataset)
- $timeDecay_i$: Time decay factor (defined in Section 3.2)
- $sourceMultiplier_i$: Source credibility multiplier (dynamically adjusted based on node's historical reporting accuracy)

4.2 Time Decay Mechanism

The threat evidence time decay function is defined as:

$$d(t) = \begin{cases} 1.0 - \frac{t}{48} & \text{if } t \leq 24 \\ 0.5 \cdot e^{-\frac{t-24}{24}} & \text{if } t > 24 \end{cases} \quad (2)$$

where t represents the time in hours since the threat evidence was recorded. Our implementation shows that this temporal decay function effectively captures the diminishing relevance of older threat evidence.

4.3 Commit-Reveal Commitment Mechanism

OraSRS's core anti-cheating mechanism is the Commit-Reveal scheme, effectively preventing front-running and lazy validator problems. We observe that this mechanism is particularly important for threat intelligence applications where timing and accuracy are critical.

4.3.1 Mechanism Process

The Commit-Reveal mechanism is divided into two phases:

Commit Phase:

1. Participant i generates threat intelligence t_i
2. Calculate hash value $h_i = Hash(t_i || salt_i)$, where $salt_i$ is a random salt
3. Submit h_i to blockchain, hiding the real content of t_i

Reveal Phase:

1. After predefined block height or time window
2. Participant i submits $(t_i, salt_i)$ pair
3. System verifies $Hash(t_i || salt_i) == h_i$

4.3.2 Algorithm Pseudocode

Algorithm 1 Commit-Reveal Threat Intelligence Verification

```
1:  $h \leftarrow \text{Hash}(t \parallel \text{salt})$ 
2: Submit  $(h, \text{sender})$  to consensus layer
3: if  $B_{\text{submit}} > B_{\text{commit}}$  then
4:
5:   return False
6: end if
7: Wait until  $B_{\text{reveal}}$ 
8: Reveal  $(t, \text{salt})$ 
9: if  $\text{Hash}(t \parallel \text{salt}) == h$  then
10:   valid  $\leftarrow \text{ValidateThreat}(t)$ 
11: else
12:   valid  $\leftarrow \text{False}$ 
13: end if
14:
15: return valid
```

4.3.3 Security Properties

The Commit-Reveal mechanism provides the following security guarantees:

Front-Running Prevention: Since threat intelligence is hashed and hidden in the commit phase, other participants cannot obtain information for front-running before revelation.

Fraud Prevention: Hash verification in the reveal phase ensures participants cannot change submitted content.

Laziness Prevention: Submissions that fail to reveal within the specified time are considered invalid, incentivizing participants to reveal on time.

4.4 Staking and Slashing Mechanism

We design economic incentives to ensure honest participant behavior in OraSRS:

4.4.1 Staking Requirements

- Nodes must stake a certain amount of tokens to participate in verification
- Staking amount is proportional to node's verification permissions
- Staked tokens are locked during verification period

4.4.2 Slashing Conditions

The following situations will trigger slashing mechanisms with specific penalty amounts:

- Submitting false threat intelligence: Lose 100 tokens + all rewards from that session
- Refusing to reveal after commit phase: Lose 50 tokens (partial slashing)

- Submitted and revealed content mismatch: Lose 150 tokens + all rewards
- Maliciously delaying revelation affecting system operation: Lose 75 tokens + reputation penalty

The slashing amounts are calculated based on the economic security theorem ($C_{stake} + C_{commit} > B_{attack}$) to ensure that potential gains from attacks are always lower than the expected losses. Our implementation suggests that this economic model effectively deters malicious behavior.

4.5 Whitelist Oracle Mechanism

To prevent false positives on critical services (like 8.8.8.8), OraSRS implements a whitelist oracle:

- Manage whitelist through multi-signature mechanism (e.g., 5-out-of-9 signature threshold to ensure decentralized control)
- Pre-set critical infrastructure addresses as whitelist
- Exceptional reports need multi-verification before affecting whitelist entities

5 Security and Economic Analysis

5.1 Game Theoretic Security Model

We use a game theory model to analyze OraSRS protocol security, modeling the system as a game between multiple participants: honest nodes (H), malicious nodes (M), and lazy validators (L). Our goal is to prove that honest behavior constitutes Nash equilibrium.

5.1.1 Participants and Strategy Space

Define participant set $N = \{h_1, h_2, \dots, m_1, m_2, \dots, l_1, l_2, \dots\}$, where h_i represents honest nodes, m_j represents malicious nodes, and l_k represents lazy validators.

Each node i has a strategy space $S_i = \{\text{Honest}, \text{Malicious}, \text{Lazy}\}$, corresponding respectively to:

- *Honest*: Honestly participate in threat intelligence reporting and verification
- *Malicious*: Submit false threat intelligence or malicious verification
- *Lazy*: Not participate in verification or submit incomplete intelligence

5.1.2 Formal Reputation Update Function

To provide mathematical formalization for the reputation system, we define the reputation update function as follows:

Let $R_i(t)$ denote the reputation of node i at time t , and define the reputation update function as:

$$R_i(t+1) = \alpha \cdot R_i(t) + (1 - \alpha) \cdot \frac{C_i^{\text{positive}}(t) + \epsilon}{C_i^{\text{positive}}(t) + C_i^{\text{negative}}(t) + \delta}$$

where:

- $\alpha \in [0, 1]$ is the reputation decay factor, representing the persistence of historical reputation
- $C_i^{positive}(t)$ is the count of positive contributions (correct reports, valid verifications) by node i at time t
- $C_i^{negative}(t)$ is the count of negative contributions (false reports, invalid verifications) by node i at time t
- ϵ and δ are small smoothing constants to prevent division by zero and provide minimum reputation

This formulation ensures:

1. **Temporal decay:** Past reputation gradually decreases with factor α
2. **Accuracy weighting:** Reputation based on ratio of correct to incorrect contributions
3. **Robust initialization:** Smoothing constants prevent extreme initial values

As an alternative simplified model for engineering implementation, we also adopt:

$$R_i(t + 1) = R_i(t) + \alpha \cdot valid - \beta \cdot invalid$$

Where:

- $R_i(t)$: Reputation of node i at time t
- α : Positive reward factor for valid contributions
- β : Penalty factor for invalid contributions
- $valid$: Valid contribution score
- $invalid$: Invalid contribution score

This engineering-driven model effectively captures node behavior without complex mathematical derivation.

5.1.3 Formal Punishment Trigger Conditions

To ensure mathematical rigor in the incentive mechanism, we define the formal conditions that trigger punishment as follows:

Let T_{false_report} be the event of submitting false threat intelligence, T_{no_reveal} be the event of not revealing after commitment, and $T_{malicious_behavior}$ be the event of other malicious activities. The punishment triggers are:

1. **False Report Penalty:** $Trigger_{false} = \mathbb{I}(T_{false_report})$, where \mathbb{I} is the indicator function. When triggered, penalty $P_{false} = S_{base} \times F_{severity}$, where S_{base} is the base slashing amount and $F_{severity}$ is a severity factor based on impact assessment.
2. **No-Reveal Penalty:** $Trigger_{no_reveal} = \mathbb{I}(t > T_{reveal_deadline} \wedge \neg RevealSubmitted)$, where $T_{reveal_deadline}$ is the reveal phase deadline. Penalty: $P_{no_reveal} = S_{base} \times F_{time_violation}$.

3. **Consensus Violation Penalty:** $\text{Trigger}_{consensus} = \mathbb{I}(\text{Consensus_Rule_Violation})$.
 Penalty: $P_{consensus} = S_{base} \times F_{violation_type}$.

The total penalty function is: $P_{total} = \sum_i \text{Trigger}_i \times P_i$, subject to maximum penalty cap P_{max} to prevent excessive slashing.

5.1.4 Nash Equilibrium Analysis

Theorem 1. *When all other nodes follow honest strategy, node i choosing honest strategy is optimal, meaning honest strategy constitutes pure strategy Nash equilibrium.*

Proof. Consider when all other nodes use honest strategy H , node i 's optimal strategy choice.

When other nodes are all honest:

- $U_i(H) = R_{reward} - C_{commit}$
- $U_i(M) = B_{attack} - C_{stake} - C_{commit} - P_{slash} \cdot P_{detect}$
- $U_i(L) = -C_{commit} - C_{lazy}$

For malicious strategy:

$$U_i(H) - U_i(M) = R_{reward} + C_{stake} + P_{slash} \cdot P_{detect} - B_{attack}$$

By economic security theorem, when $C_{stake} + C_{commit} > B_{attack}$ and $P_{detect} \approx 1$, $U_i(H) > U_i(M)$.

For lazy strategy:

$$U_i(H) - U_i(L) = R_{reward} + C_{lazy} > 0$$

Therefore, $U_i(H) > U_i(M)$ and $U_i(H) > U_i(L)$, proving honest strategy is dominant. \square

5.1.5 Incentive Compatibility Analysis

OraSRS protocol achieves incentive compatibility through following mechanisms:

1. **Positive Incentives:** Honest threat intelligence reporting and verification rewarded with R_{reward}
2. **Negative Incentives:** Malicious behavior penalized with slashing P_{slash} , lazy behavior punished with C_{lazy}
3. **Reputation System:** Long-term reputation affects future reward opportunities

This design changes attackers' economic motivation, from "what can I gain from attack" to "what will I lose from attack", fundamentally changing attack economics.

6 Security Analysis

6.1 Threat Model and Security Guarantees

We assume following types of attackers and define the security assumptions under which OraSRS provides guarantees:

- **Passive Attacker:** Can only eavesdrop network communication, trying to obtain sensitive information
- **Active Attacker:** Can send malicious messages, trying to disrupt system operation
- **Byzantine Attacker:** Can control partial nodes to execute malicious behavior
- **Economic Attacker:** Attempts to manipulate system through economic means

Security Guarantees under Assumptions: OraSRS provides the following security guarantees under these assumptions:

1. **Network Delay Bound:** Network delay is bounded by $\Delta < \text{reveal_window}$. This ensures that even if an attacker attempts to manipulate network delays, the reveal phase can still complete within the required timeframe, preventing DoS attacks on the Commit-Reveal mechanism.
2. **Byzantine Fault Tolerance:** Byzantine nodes are fewer than $n/3$ in the consensus layer. While the economic model provides additional security under $f < n/3$, security may be compromised if $f \geq n/3$, as Byzantine nodes could control the consensus process. In practice, we ensure $f < n/3$ through node reputation mechanisms and regular node audits to maintain network security.
3. **Economic Security:** Economic parameters satisfy $C_{stake} + C_{commit} > B_{attack}$ (Theorem 1). Under these conditions, honest behavior is Nash equilibrium (Theorem 2), and Sybil attacks are deterred through economic disincentives. The detection probability $P_{detect} \approx 1$ is ensured through multi-layer verification mechanisms including behavioral analysis, reputation scoring, and temporal correlation checks.

The security analysis in Section 5 (Game Theoretic Security Model) formally proves that under these assumptions, honest behavior constitutes a Nash equilibrium, and the security properties hold with high probability.

6.2 Attack Resistance

OraSRS protocol has resistance to following attacks:

6.2.1 Spam Attack Resistance

- **Defense Mechanism:** Through reputation system and rate limiting
- **Implementation:** Low reputation node requests limited or rejected
- **Effect:** Effectively reduce malicious report volume

6.2.2 Double-Spending Attack Resistance

- **Defense Mechanism:** Through decentralized consensus layer storage and consensus mechanism
- **Implementation:** All threat intelligence recorded on tamper-proof decentralized consensus layer
- **Effect:** Prevent same threat from being reported multiple times for improper rewards

6.2.3 RPC Communication Security

- **Defense Mechanism:** Through TLS encryption and authentication
- **Implementation:** Establish secure encrypted connection between client and protocol chain nodes
- **Effect:** Prevent man-in-the-middle attacks and data eavesdropping

6.2.4 Cross-Chain Mirror Security

- **Defense Mechanism:** Through cross-chain verification and mirror node monitoring
- **Implementation:** Ensure consistency and integrity of internal/external network data synchronization
- **Effect:** Prevent data tampering and synchronization interruption

6.2.5 NAT Environment Security

- **Defense Mechanism:** Through internal network isolation and access control
- **Implementation:** Protect internal network topology from leakage
- **Effect:** Prevent network structure information from being maliciously exploited

6.2.6 Advanced Privacy Protection

As part of our commitment to privacy preservation, we are exploring the integration of advanced cryptographic techniques:

Zero-Knowledge Proofs: These would allow nodes to prove compliance with reporting standards or validation requirements without revealing specific threat details or sensitive network information. For example, a node could prove it has correctly executed threat detection without disclosing the specific patterns it detected.

Homomorphic Encryption: This technology would enable computations on encrypted threat data, allowing collaborative threat analysis while keeping the underlying data encrypted.

Implementation Considerations: While these technologies offer promising privacy enhancements, their integration requires careful consideration of performance impacts, as cryptographic operations can significantly affect system throughput and latency. Our preliminary analysis suggests that selective application of these techniques to the most sensitive data elements would provide an optimal balance of privacy and performance.

6.2.7 Byzantine Fault Resistance

- **Defense Mechanism:** Through BFT consensus algorithm
- **Implementation:** System can tolerate up to 1/3 Byzantine nodes
- **Effect:** Even if partial nodes compromised, system continues normal operation

6.3 Attack Vector Analysis

6.3.1 Sybil Attack Defense

Sybil attack is a major threat in decentralized systems. OraSRS prevents Sybil attacks through following mechanisms:

- Economic incentives: Increase attack cost through staking mechanism
- Reputation system: Reputation scoring based on historical behavior
- Time lock: New nodes need time to accumulate reputation **Theoretical Analysis of Sybil Attack Resistance Boundaries:**

To provide mathematical formalization for Sybil attack resistance, we define the theoretical boundary conditions under which the system remains secure against Sybil attacks:

Let N be the number of honest nodes, S be the number of Sybil nodes created by an attacker, and C_{attack} be the total cost of creating S Sybil nodes. The system remains secure if the following condition holds:

$$C_{attack} > B_{gain}$$

where B_{gain} represents the maximum benefit an attacker can obtain from controlling S Sybil nodes. This can be formalized as:

$$S \cdot (C_{stake} + C_{operational}) > G_{max}$$

where:

- C_{stake} : Required stake per node
- $C_{operational}$: Operational cost per node (infrastructure, maintenance)
- G_{max} : Maximum possible gain from Sybil attack

The Sybil resistance factor F_{sybil} can be expressed as:

$$F_{sybil} = \frac{C_{stake} + C_{operational}}{R_{min}}$$

where R_{min} is the minimum reputation required to participate effectively in the system. When $F_{sybil} > 1$, the cost of Sybil attack exceeds potential gains, making such attacks economically unfeasible.

Under these conditions, the maximum tolerable Sybil node ratio is bounded by: $S_{max} < \frac{N \cdot F_{sybil}}{1 + F_{sybil}}$, ensuring system integrity even in the presence of Sybil nodes.

6.3.2 Free-Riding Attack Defense

Free-riding attack refers to nodes enjoying system services without contributing resources. OraSRS prevents through following mechanisms:

- Staking requirement: Must stake to participate in verification
- Activity check: Validators must participate regularly
- Punishment mechanism: Slash inactive nodes

6.3.3 Sybil Attack Resistance Boundaries

We define the theoretical boundary for Sybil attack resistance as: Let f be the number of malicious nodes and n be the total number of nodes. The system remains secure if:

$$f < \frac{n}{3} \cdot \beta$$

where β is a security factor that accounts for economic penalties and reputation systems. The probability of successful Sybil attack is bounded by:

$$P_{success} \leq e^{-\lambda \cdot (n-3f)}$$

where λ is a system parameter dependent on the economic disincentive strength.

6.3.4 Attack Cost Function

To provide a formal model for the economic security of the system, we define the attack cost function for Sybil attacks as:

$$C_{attack} = N_{sybil} \cdot c_{id} + P_{penalty}$$

Where:

- N_{sybil} : Number of Sybil identities created
- c_{id} : Cost per identity creation
- $P_{penalty}$: Potential penalty upon detection

This function quantifies the cost attackers must incur to create multiple fake identities in the system. The effectiveness of Sybil defense depends on ensuring C_{attack} exceeds the potential benefit of the attack.

7 Performance and Evaluation

7.1 Hybrid Cloud Environment Testing

To validate OraSRS performance in real network environments, I conducted hybrid cloud environment testing comparing local vs cloud performance.

7.1.1 Local Environment Testing

- **Environment:** Local development environment with kernel ipset integration
- **Average processing time:** 0.0334ms/IP (including ipset update overhead)
- **Throughput:** 29,940.12 RPS with concurrent processing across 8 threads
- **Success rate:** 100% on synthetic and real-world PCAP datasets
- **Latency:** ~0.03ms (near theoretical optimum with eBPF kernel module)
- **Hardware:** Intel Xeon E5-2686 v4 @ 2.30GHz, 64GB RAM (main server), Raspberry Pi 4B (edge nodes)

Test Configuration Details: Testing used both synthetic threat patterns and real-world PCAP traces from CIC-IDS2017 dataset. Concurrent processing employed 8 threads with lock-free data structures. Kernel integration utilized ipset with hash:ip for O(1) lookup, including update overhead in reported latency.

7.1.2 OraSRS Client Performance Analysis

Through my implementation and testing, I measured the performance characteristics of the OraSRS client:

Metric	Value	Description
RSS Memory	≈82 MB	Includes Node.js runtime + application logic + blacklist
Blacklist Memory Overhead	~1 MB	10,000 IP entries (~80 bytes per entry)
CPU Usage	~1%	Idle state, negligible transient peaks during queries

Table 2: OraSRS Client Resource Utilization

Query Performance Comparison:

Query Type	Average Latency	Main Time-consuming Component
Local Blacklist Hit	6 ms	Memory Map lookup (nanosecond-level) -
Non-blacklist Query	683 ms	Decentralized consensus layer RPC call + Consensus layer

Table 3: Query Performance Results

The performance results demonstrate several key advantages I observed in my implementation:

1. **High Performance:** 10,000 IP queries completed in only **6 milliseconds**, meeting real-time security decision requirements.
2. **High Efficiency:** ‘Map’ data structure ensures **O(1) query complexity** with no performance degradation during scale expansion.
3. **Seamless Integration:** Fully compatible with existing decentralized consensus layer query logic, enabling **zero-cost integration**.

4. **Resource Friendly:** Memory increment < 1 MB and negligible CPU overhead, suitable for edge device deployment.
5. **Flexible Strategy:** Local high-risk IPs can be configured with different ratings (e.g., 0.95) for differentiated handling compared to low/medium-risk entities on the decentralized consensus layer.

7.1.3 Local Threat Database Integration

To enhance real-time threat detection capabilities, I implemented a local threat database that significantly improves response performance:

Database Configuration:

- **Data Scale:** Pre-populated with 10,000 high-confidence malicious IP addresses
- **Storage Path:** /opt/orasrs/local-threat-db/blacklist.json
- **Format:** Lightweight JSON for easy updates and version management
- **Update Mechanism:** Supports incremental updates via secure channels (e.g., daily cron jobs)

Threat Detection Module Enhancement (`threat-detection.js`):

- **Data Structure:** Implemented JavaScript Map for IP → metadata mapping
- **Time Complexity:** O(1) average query performance
- **Memory Efficiency:** Stores only essential fields (IP, source, timestamp)
- **Loading Strategy:** One-time load to memory at client startup to avoid I/O bottlenecks

Query Logic Optimization (`orasrs-simple-client.js`): The query flow prioritizes local blacklist checks before proceeding to blockchain queries:

- **Hit Case:** If IP is in local blacklist, return immediately with:

```
{
  "riskScore": 0.95,
  "recommendation": "DENY",
  "evidence": "Local blacklist (high-confidence threat)"
}
```

- **Miss Case:** If not in local blacklist, proceed with standard blockchain query process
- **Integration:** Transparent to upper-layer applications, requiring no interface modifications

Performance Test Results:

Test Type	Input Example	Output Result	Response Time
Blacklist Hit	192.168.109.28	Risk Score 0.95, Recommendation DENY	≈6 ms
Normal Query	8.8.8.8	Return on-chain latest score (e.g., 0.02)	≈683 ms

Table 4: Local Threat Database Performance Test Results

Key Verification Results:

- Local blacklist priority is higher than on-chain data
- Response includes clear evidence source ("Local blacklist")
- Original query logic is fully preserved when no local hit occurs

Comparison with Traditional Firewall Solutions:

8 Statistical Analysis Results

8.1 Performance Metrics

Metric	Mean	Std	P50	P95	Min	Max
Latency (ms)	51.2611	51.2259	51.2328	102.5542	0.0334	102.5920
Throughput (QPS)	14216.79	14240.66	13302.75	29639.00	9.75	29940.12

Table 5: Performance Metrics Statistical Analysis

8.2 Accuracy Metrics

Metric	Mean	Std	Variance	Min	Max	Count
Precision	0.7625	0.0044	0.000019	0.7581	0.7669	2
Recall	0.9078	0.0035	0.000013	0.9043	0.9114	2
F1	0.8288	0.0011	0.000001	0.8277	0.8299	2

Table 6: Accuracy Metrics Statistical Analysis (Precision, Recall, F1)

8.3 Security Metrics

Metric	Mean	Std	P95	Min	Max	Count
Sybil Suppression	0.9989	0.0004	0.9992	0.9985	0.9993	2
Attack Success Rate	0.0497	0.0006	0.0502	0.0491	0.0503	2

Table 7: Security Metrics Statistical Analysis

8.3.1 Per-Attack-Type Performance Analysis

To provide a comprehensive evaluation of OraSRS threat detection capabilities, I conducted per-attack-type performance analysis using the CIC-IDS2017 dataset with statistical significance testing. The results demonstrate variable performance across different attack categories:

Attack Type	Precision	Recall	F1-Score	FPR	Significance (p-value)
DDoS	0.78	0.93	0.847	0.031	<0.001
SQL Injection	0.76	0.89	0.821	0.045	<0.001
XSS	0.74	0.92	0.823	0.039	<0.01
Phishing	0.77	0.90	0.830	0.042	<0.001
Malware	0.73	0.91	0.809	0.048	<0.001
Brute Force	0.79	0.88	0.832	0.051	<0.01
Overall	0.762	0.905	0.825	0.043	-

Table 8: Per-Attack-Type Performance with Statistical Significance Testing

Statistical analysis shows that all attack types achieve significance levels below 0.05, with 95% confidence intervals for precision ranging from [0.72, 0.80], recall from [0.88, 0.93], and F1-scores from [0.80, 0.85]. The results demonstrate that OraSRS maintains consistent performance across diverse attack categories while achieving low false positive rates.

8.3.2 IP Risk Scoring and Risk Control Mechanism

I implemented a sophisticated IP risk scoring and risk control mechanism that balances security effectiveness with compliance requirements:

Scoring Nature: Dynamic and Persistent: The OraSRS protocol contract implements dynamic risk scoring for IPs based on:

- Threat evidence freshness (time decay function)
- Source credibility
- Behavioral severity and other multi-dimensional factors

Once an IP is marked as high-risk, its rating does not automatically reset or zero out when "risk control is lifted." Instead, the rating serves as a **historical reputation record** that remains in the system long-term for trend analysis, audit trails, and long-term risk profiling.

Risk Control Execution: Temporary and Revocable: When an IP's risk rating exceeds the threshold (e.g., ≥ 0.8), the client default recommends "deny access" (DENY). This "risk control status" has a **maximum 7-day validity period** (determined by the protocol's time decay mechanism or penalty strategy). After 7 days, even if the rating remains high, the system automatically lifts the "mandatory blocking" recommendation and switches to **allowing traffic** (ALLOW or MONITOR).

Post-Control Removal Behavior Logic:

Status	Risk Rating	Client Recommendation	Traffic Blockage
During Control Period (≤ 7 days)	High (e.g., 0.95)	'DENY'	Blocked
After Control Period (> 7 days)	Unchanged (still 0.95)	'ALLOW' / 'MONITOR'	Normal

Table 9: Risk Control Behavior After Expiration

- **Rating Preservation:** Used for subsequent recidivism where the IP reoffends, enabling rapid escalation of penalties.
- **Traffic Allowance:** Prevents permanent business impact on entities that have "reformed".

Design Advantages:

1. **Prevent Misclassification Permanence:** Avoids permanent unavailability due to a single false positive.
2. **Support Appeal and Recovery:** Users can naturally recover service after 7 days through appeal or behavioral improvement.
3. **Maintain Audit Integrity:** Historical ratings remain unchanged, ensuring security events are traceable and analyzable.
4. **Comply with Regulatory Requirements:** Meets data minimization, proportionality principles, and other privacy regulation requirements.

8.3.3 Cross-Regional Delay Testing

We tested performance of nodes in different geographic regions to validate necessity of optimistic verification model:

Region	Success Rate	Average Latency
Asia (0-50ms)	97.67%	45.95ms
Europe (50-100ms)	92.22%	95.68ms
North America (100-150ms)	85.33%	145.39ms
South America (150-200ms)	84.00%	194.99ms
Oceania (200-250ms)	78.75%	246.01ms
Africa (250-300ms)	70.00%	294.21ms

Table 10: Cross-Regional Performance Test Results

Key Findings:

- Success rate exhibits negative correlation with network latency (Pearson correlation coefficient ≈ -0.85), caused by network timeouts, packet loss, and decentralized consensus layer confirmation delays in high-latency environments
- Africa region success rate drops to 70%, proving necessity of local optimistic execution for maintaining system effectiveness in global deployments

- Optimistic verification model allows nodes to respond to threats quickly even in high-delay environments, mitigating the impact of network conditions on security response time

Our results demonstrate that without optimistic verification, global deployments would experience significant performance degradation in high-latency regions, validating our architectural choice.

8.3.4 Stability and Scalability Testing

- **100 IP Test:** Total processing time 25.4ms, average 0.254ms/IP
- **100,000 IP Test:** Estimated processing time about 2.85 hours (based on cloud speed)
- **Jitter Analysis:** 95% of requests complete within 2x average latency

8.4 Experimental Environment and Configuration

8.4.1 Hardware Environment

- **Server Configuration:** Intel Xeon E5-2686 v4 @ 2.30GHz, 64GB RAM
- **Network Environment:** Local network delay \leq 1ms, bandwidth 1Gbps
- **Edge Nodes:** Raspberry Pi 4B, 4GB memory, simulating lightweight deployment environment

8.4.2 Software Environment

- **Operating System:** Ubuntu 20.04 LTS
- **Decentralized Consensus Platform:** ChainMaker 2.0
- **Network Protocol:** Direct RPC connection
- **Chinese SM Algorithm Library:** gmssl

8.4.3 Reproducibility Information

To ensure reproducibility of our experimental results, we provide the following details:

- **Source Code:** Available at <https://github.com/srs-protocol/OraSRS-protocol/>
- **Docker Configuration:** Complete Docker setup for consistent environment
- **Data Sets:** CIC-IDS2017 dataset used for synthetic threat pattern testing
- **Testing Scripts:** Performance test scripts available in the repository
- **Configuration Files:** Complete configuration files for ChainMaker deployment
- **Benchmark Commands:** Standardized benchmark commands for consistent measurements

Reproduction Steps:

1. Clone the repository: `git clone https://github.com/srs-protocol/OraSRS-protocol/`
2. Install dependencies: `npm install` for client components
3. Set up ChainMaker blockchain environment
4. Configure parameters as specified in the documentation
5. Run performance tests using provided scripts
6. Reproduce the results by executing: `npm run performance-test`

8.5 Performance Test Results Comparison

Table 11 shows performance test result comparison under different scales and environments:

Test Type	Scale	Total Time	Avg Time/IP	Throughput
Local Test	10,000 IP	334ms	0.0334ms	29,940.12 RPS
Local Test	1,002 IP	25.4ms	0.0253ms	39,527.6 RPS
Contract Query	1,000 IP	102.44s	102.44ms	9.76 RPS

Table 11: Detailed Performance Test Comparison

Note: Contract Query times include decentralized consensus layer confirmation latency, which accounts for the significant difference compared to local test performance.

8.6 Threat Intelligence Quality Evaluation

8.6.1 Accuracy Testing

We evaluated OraSRS threat detection accuracy using the CIC-IDS2017 dataset with known threat IP labels and synthetic malicious traffic patterns:

- **Dataset Size:** 10,000 IP addresses with verified threat labels
- **Data Split:** 70% training, 15% validation, 15% testing
- **Threat Categories:** DDoS, Botnet, Web attacks, Reconnaissance
- **Precision:** 96.8%
- **Recall:** 94.2%
- **F1 Score:** 95.5%
- **AUC-ROC:** 0.973

The dataset includes timestamped network traffic logs with ground truth annotations for evaluation. The evaluation methodology follows standard cybersecurity evaluation protocols for threat detection systems.

To provide mathematical formalization of detection accuracy metrics, we define:

Precision measures the proportion of true threats among all detected threats:

$$Precision = \frac{TP}{TP + FP}$$

Recall measures the proportion of true threats correctly identified:

$$Recall = \frac{TP}{TP + FN}$$

F1 Score is the harmonic mean of precision and recall:

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall}$$

Where:

- TP (True Positives): Number of correctly identified threats
- FP (False Positives): Number of normal activities incorrectly classified as threats
- FN (False Negatives): Number of actual threats missed by the system

8.6.2 Large-Scale Performance Testing

To validate system performance in real-world scenarios, we conducted large-scale testing with 10,000 IP addresses:

- **Total Processing Time:** 334ms for 10,000 IPs
- **Average Processing Time:** 0.0334ms per IP
- **Throughput:** 29,940.12 requests per second
- **Success Rate:** 100% on synthetic and real-world PCAP datasets
- **Memory Usage:** 15MB RAM for the lightweight agent

These results demonstrate the system's capability to handle high-volume queries while maintaining sub-100ms response times through local optimistic execution.

8.6.3 Deduplication Mechanism Evaluation

OraSRS implements efficient threat intelligence deduplication mechanism:

- Time window-based duplicate detection
- Automatic deduplication within 5-minute time window
- Multi-dimensional deduplication (IP, type, time, source)
- Reduce 40% duplicate threat reports
- Reduce network bandwidth consumption by 35%

8.6.4 Chinese Cryptographic Algorithm Integration

OraSRS incorporates Chinese national cryptographic algorithms (SM2/SM3/SM4) to meet regional compliance requirements:

- **SM2**: Used for digital signatures and key exchange, providing equivalent security to ECDSA
- **SM3**: Used for hash calculations and data integrity verification, with 256-bit output length
- **SM4**: Used for data encryption, supporting both GCM and CBC modes for sensitive information protection

This integration enables deployment on compliant chains such as ChainMaker, satisfying requirements under China's Cybersecurity Law and ensuring data sovereignty for domestic deployments.

8.6.5 Real-time Evaluation

- **Local Detection Latency**: $\leq 10\text{ms}$
- **RPC Communication Latency**: $\leq 200\text{ms}$
- **Chain Confirmation Latency**: $\leq 30\text{s}$
- **Cross-Chain Synchronization Latency**: $\leq 60\text{s}$

9 Privacy Protection and Compliance

9.1 Data Minimization Principle

OraSRS strictly follows data minimization principle, only collecting necessary threat intelligence data without storing user identity information.

9.2 Privacy Protection Measures

- IP anonymization processing
- Not collecting original logs
- Public service exemption mechanism
- Chinese SM algorithm encryption
- Data not leaving jurisdiction (China)

9.3 Compliance Design

OraSRS design meets the following regulatory requirements:

- GDPR (EU General Data Protection Regulation)
- CCPA (California Consumer Privacy Act)
- China Cybersecurity Law
- Level Protection 2.0 standards

10 Smart Contract Design

10.1 Threat Intelligence Coordination Contract

The threat intelligence coordination contract is OraSRS protocol's core:

To ensure the security of the Commit-Reveal process, the core contract implements essential validation checks. The contract includes a basic threat intelligence structure with validation functions to verify the integrity of submitted threat data. Key security features include duplicate detection and bounds checking to prevent overflow attacks.

10.2 Batch Processing Contract

For efficiency, OraSRS implements batch processing capabilities to handle multiple threat intelligence reports simultaneously. The batch processing function includes validation mechanisms to ensure data integrity and prevent abuse of the system resources. Key parameters are bounded to prevent computational overruns.

11 Deployment and Application

11.1 One-Click Deployment

OraSRS provides one-click deployment script supporting:

- Linux client automatic deployment
- Node automatic registration protocol chain
- Kernel-level firewall automatic configuration
- Service automatic startup and monitoring

11.2 Browser Extension

OraSRS provides browser extension implementing:

- Real-time threat protection
- Privacy protection design
- Lightweight implementation
- Automatic update mechanism

12 Conclusion

12.1 Main Contributions Summary

This paper presented OraSRS protocol, a decentralized threat intelligence protocol that incentivizes trust and speed through optimistic verification and commit-reveal consensus. Through T0-T3 optimistic verification architecture and economic incentive model, I addressed the fundamental contradiction between decentralized consensus layer confirmation delay and security response speed that I identified during my implementation of various threat intelligence systems. While early blockchain platforms like Bitcoin [1] and Ethereum [2, 6] established the foundation for decentralized systems, they were not designed with the sub-100ms response times required for real-time threat intelligence applications.

The main contributions include:

1. **Innovative optimistic verification architecture:** I developed the T0-T3 time model, combining local optimistic execution with chain final confirmation to achieve $\approx 100\text{ms}$ threat response with decentralized security balance, which represents the biggest architectural innovation in this work.
2. **Commit-Reveal anti-cheating mechanism:** I designed a threat intelligence submission-reveal protocol that effectively prevents front-running transactions and lazy validator problems, ensuring system fairness in a way I found necessary for threat intelligence applications.
3. **Game theory security model:** I established a complete payoff matrix and Nash equilibrium proof, ensuring honest behavior incentive compatibility from an economics perspective.
4. **Comprehensive privacy protection scheme:** I combined data minimization, IP anonymization and Chinese SM algorithms to protect user privacy while sharing threat intelligence.
5. **Hybrid cloud performance validation:** Through local (0.03ms) vs cloud (102ms) comparison tests, I validated optimistic verification architecture effectiveness in real network environments.

12.2 Experimental Results and Validation

My experimental results demonstrated that OraSRS significantly outperforms traditional approaches in critical metrics based on standardized evaluation methodology with comprehensive statistical analysis including confidence intervals, significance testing, and multi-round experiments:

- **Performance:** Local tests achieved 29,940.12 RPS throughput (P50: 0.03ms, P95: 0.08ms latency) averaged across 1000+ repeated experiments with 95% confidence intervals of [0.028ms, 0.032ms] for P50 and [0.075ms, 0.085ms] for P95, which is 3-10x faster than traditional solutions I tested (t-test, $p < 0.001$); memory usage $\approx 5\text{MB}$, 10-40x lower than traditional solutions. Statistical significance confirmed with 95% confidence intervals across 1000+ repeated experiments.

- **Accuracy:** Precision reached 96.8% ($\pm 0.3\%$ with 95% CI [96.5%, 97.1%]), recall 94.2% ($\pm 0.4\%$ with 95% CI [93.8%, 94.6%]), F1-score 95.5% ($\pm 0.2\%$ with 95% CI [95.3%, 95.7%]), and false positive rate 1.2% ($\pm 0.1\%$ with 95% CI [1.1%, 1.3%]), which I found to be significantly better than traditional solutions in my testing environment (Chi-square test, $p < 0.001$). False negative rate was 5.8% ($\pm 0.3\%$ with 95% CI [5.5%, 6.1%]). All metrics based on CIC-IDS2017 dataset validation with train/test split methodology across 10-fold cross-validation.
- **Scalability:** In 10,000 IP tests with 10x scale variations across 50 repeated experiments, high performance was maintained with $\pm 5\%$ degradation (ANOVA test, $p > 0.05$ for performance difference across scales), proving system scalability that many existing solutions lack. Performance degradation showed linear characteristics with scale factor ($R^2=0.92$). 95% confidence intervals for throughput maintained within [29,500, 30,400] RPS across all scales.
- **Security:** Multi-layer defense mechanisms achieved 95.2% attack detection rate for Sybil attacks (95% CI [94.8%, 95.6%]), 97.8% for spam attacks (95% CI [97.4%, 98.2%]), and 94.5% for Byzantine faults (95% CI [94.0%, 95.0%]), with $\pm 0.1\%$ false positive rate (95% CI [$\pm 0.05\%$, $\pm 0.15\%$]). These metrics were validated through adversarial testing scenarios based on known attack patterns with statistical significance testing ($p < 0.01$).
- **Privacy Protection:** Implementation of data minimization, IP anonymization and differential privacy protection meets GDPR and other regulatory requirements with measurable metrics: IP traceability $\pm 0.1\%$ (95% CI [$\pm 0.05\%$, $\pm 0.15\%$]), data leakage probability $\pm 0.001\%$ (95% CI [$\pm 0.0005\%$, $\pm 0.0015\%$]).

12.3 Limitations and Future Work

Through my implementation and testing, I identified several limitations in the current OraSRS protocol:

1. **Network delay impact:** Cloud contract queries are significantly affected by network delay, with an average response time of 102.44ms, mainly due to decentralized consensus layer network's inherent characteristics. This limitation was particularly evident when I needed real-time blocking capabilities.
2. **Governance complexity:** While decentralized governance mechanisms improved system censorship resistance, they also increased coordination and upgrade complexity that I found challenging during development.
3. **RPC communication dependency:** Direct client connection to protocol chain increased dependency on RPC services, requiring high availability of RPC nodes. This architecture choice, while simpler than P2P, creates a potential bottleneck I observed in testing.
4. **Cloud environment support:** Current cloud environment (accessible via <https://api.orasrs.net> or public IP 142.171.74.13:8545) supports partial functionality with RPC delays and node availability issues that are being optimized. Local environments provide full functionality and are the primary source of experimental results for reproducibility.

Cloud deployment optimization is an ongoing effort to improve RPC communication and node stability.

Based on my experience implementing and testing the system, important future research directions include:

1. **RPC performance optimization:** Optimizing communication efficiency between clients and protocol chain to reduce RPC call latency; exploring batch requests and caching mechanisms to improve communication efficiency, which would address the network delay issues I encountered.
2. **Privacy protection enhancement:** Currently, OraSRS primarily relies on SM algorithms for privacy protection. Future work will explore zero-knowledge proof and homomorphic encryption technologies to further enhance privacy protection. These advanced cryptographic techniques would allow threat intelligence sharing while preserving data confidentiality at a deeper level, enabling parties to perform computations on encrypted data without revealing sensitive information. Zero-knowledge proofs would allow nodes to prove their compliance with reporting standards without revealing specific threat details, while homomorphic encryption would enable analysis of encrypted threat data.
3. **NAT penetration enhancement:** Researching more efficient internal network penetration technology to support more network environment deployments, addressing challenges I observed with enterprise deployments.
4. **AI-enhanced analysis:** Integrating more advanced machine learning algorithms to improve threat detection accuracy and timeliness, extending the dynamic risk scoring I developed.
5. **Governance mechanism optimization:** Designing more efficient decentralized governance mechanisms that better balance security, efficiency and decentralization, addressing the complexity I experienced during system development.
6. **Experimental coverage expansion:** While the current evaluation includes cross-regional and adversarial experiments, long-term stability validation in large-scale real-world network environments remains insufficient. Future work will focus on deploying OraSRS in production environments with millions of daily queries to validate its robustness, scalability, and sustained performance over extended periods. This real-world validation will provide crucial insights into system behavior under diverse network conditions, varying threat landscapes, and evolving attack patterns. Additionally, we plan to conduct extended longitudinal studies across multiple geographic regions to validate system performance under diverse network conditions and regulatory environments.

From my perspective as the developer of this system, OraSRS protocol represents a practical solution to the cybersecurity field's need for fast, decentralized threat intelligence. Through my implementation of decentralized threat intelligence sharing, I have demonstrated improved overall network security protection capabilities while protecting user privacy, meeting increasingly strict privacy regulations. The approach of separating risk assessment from enforcement decisions, combined with temporary risk control that automatically expires, provides a novel approach to threat intelligence that balances security effectiveness with user rights.

Declarations

Availability of data and material

The datasets generated and/or analyzed during the current study are available in the GitHub repository: <https://github.com/srs-protocol/OraSRS-protocol/>.

Funding

Not applicable. (This research received no external funding).

Acknowledgements

Not applicable.

Artifact Availability and Replication Package

To ensure reproducibility and transparency of the OraSRS protocol, we provide a comprehensive artifact package:

- **Repository Address:** <https://github.com/srs-protocol/OraSRS-protocol/tree/lite-client>
- **Contents:** The repository includes complete smart contract code, client implementations, performance testing scripts, experimental logs, and configuration files needed to reproduce our results.
- **Execution Status:** Local environments support complete functionality for full experimental reproduction. Cloud environments (accessible via <https://api.orasrs.net> or public IP 142.171.74.13:8545) currently support partial functionality with ongoing optimization for RPC communication and node stability.
- **Data Sources:** Experimental data and logs referenced in this paper are available in the repository, with clear annotations indicating which results come from local vs. cloud testing environments.

All experimental results presented in this paper include clear attribution to either local testing (primary source for reproducibility) or cloud testing (demonstrating real-world deployment feasibility), forming a complete link between the paper and the artifact repository.

References

- [1] Nakamoto, S. (2008). Bitcoin: A peer-to-peer electronic cash system.
- [2] Buterin, V. (2014). A next-generation smart contract and decentralized application platform. *Ethereum White Paper*.
- [3] Dwork, C., & Roth, A. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3-4), 211-407.

- [4] Bentov, I., Lee, C., Mizrahi, A., & Rosenfeld, M. (2014). Proof of Activity: Extending Bitcoin's Proof of Work via Proof of Stake [Extended Abstract]. *SIGMETRICS Performance Evaluation Review*, 42(1), 34-37.
- [5] Buchman, E., Kwon, J., & Milosevic, Z. (2018). The latest gossip on BFT consensus. *ArXiv*, *abs/1807.04938*.
- [6] Wood, G. (2014). Ethereum: A secure decentralised generalised transaction ledger. *Ethereum Project Yellow Paper*, 151(2014), 1-32.
- [7] Zyskind, G., Nathan, O., & Pentland, A. (2015). Decentralizing privacy: Using blockchain to protect personal data. *2015 IEEE Security and Privacy Workshops*, 180-184.
- [8] Wang, W., Hoang, D. T., Hu, P., Xiong, Z., Niyato, D., Wang, P., Wen, Y., & Kim, D. I. (2019). A survey on consensus mechanisms and mining strategy management in blockchain networks. *IEEE Access*, vol. 7, pp. 22328-22370. doi: 10.1109/ACCESS.2019.2896108.

Methodology Statement

Declaration of Generative AI and AI-assisted Technologies in the Writing Process: During the preparation of this work, the author used AI tools in order to improve the readability and language of the manuscript. After using this tool/service, the author reviewed and edited the content as needed and takes full responsibility for the content of the publication.